

Raytracing Project

Jared Givens, Ahbi Sohal, Noah Krim

28 May, 2023

Description

The program outputs color values of pixels with raytracing techniques. It begins by initializing virtual geometry and a camera. Then it sends rays from the camera into the virtual space.

For each pixel a series of uniformly random rays are cast to gather a color value. the color values are then averaged to achieve anti aliasing.

Each ray iterates over the scene's geometry to determine the nearest intersection. The nearest intersection's color is then multiplied against the ray's running total color value. The ray then reflects and scatters or refracts based on the properties of the material that was intersected. The process of accumulating color and bouncing continues until the ray fails to collide with an object or reaches the maximum ray depth. If the ray fails to collide, the program adds the contribution of the skybox to complete the color. If the ray reaches the maximum bounce depth the color is set to black.

Justification

Our serial prototype was only capable of outputting a single frame in 5.971 seconds (wall-clock time 2020 Macbook 2 GHz 16 GB). This program is an excellent candidate for CUDA, because the code of each ray can run in parallel. One of the biggest difficulties is partitioning the geometry to be intersected. The majority of the time is spent computing the ray intersections and bounces which have no data dependencies. Additionally the majority of code is linear algebra with a few branches that could be removed with better equations.

Qualification

Our team can complete this project because we completed a [serial prototype](#). We followed the tutorial [Ray Tracing In OneWeekend](#) the program writes a ppm to stdout which can be converted to a PNG. If time permits we have plans to try to integrate OpenGL BVH, triangles, model loading, lights and textures. The project uses CMake to create the makefile. to create the makefile.

Here is a frame from the serial output:

